

Name: Nachiketh Reddy

ID: 2117731

**QUESTION 1:**

Bloom Filter Consider a Bloom filter with 16 slots that uses 2 hash functions. To define the index position to set in the array, we will use the following procedure: Generate a hash value using SHA256 on your depaulid and a number. For example, to generate hash value of number x, you will need to run: Use the bloomfilter.py file shared in the class: (a) Add the strings '1', '3', '5', ..., '19' to this Bloom filter and write down which bits have been set. (b) Test the resulting Bloom filter for the values '2', '4', '6', ..., '20'. Do you get any false positives, and if so, what are they?

**QUESTION 1**

```
In [10]: import hashlib

class BloomFilter:
    def __init__(self, size, hash_functions):
        self.size = size # Size of the filter (number of bits)
        self.hash_functions = hash_functions
        self.bit_array = [False] * size

    def add(self, element):
        for i in range(self.hash_functions):
            index = self._hash(element, i)
            self.bit_array[index] = True

    def contains(self, element):
        for i in range(self.hash_functions):
            index = self._hash(element, i)
            if not self.bit_array[index]:
                return False
        return True

    def _hash(self, element, i):
        # Generate hash value using SHA256 using DePaul ID and element
        hash_value = hashlib.sha256(bytes(f"npamamah {element}", encoding='utf-8')).hexdigest()
        if i == 0:
            hash_function = int(hash_value[0], 16)
        else:
            hash_function = int(hash_value[1], 16)
        # Return the hash value modulo the size of the filter
        return hash_function % self.size

# Creating a Bloom filter with 16 slots and 2 hash functions
filter_size = 16
num_hash_functions = 2
obj_bloom = BloomFilter(filter_size, num_hash_functions)

# Add the elements '1', '3', '5', ..., '19' to the Bloom filter
for i in range(1, 20, 2):
    obj_bloom.add(str(i))
    print(f"Bit at index {obj_bloom._hash(str(i), 0)} and {obj_bloom._hash(str(i), 1)} is set for {i}")

# Test the presence of elements '2', '4', '6', ..., '20' in the Bloom Filter
values_to_check = list(range(2, 20, 2))
for value in values_to_check:
    if obj_bloom.contains(str(value)):
        print(f"{value} may be in the set.")
    else:
        print(f"{value} is definitely not in the set.")
```

```
Bit at index 4 and 8 is set for 1
Bit at index 7 and 2 is set for 3
Bit at index 15 and 1 is set for 5
Bit at index 4 and 8 is set for 7
Bit at index 10 and 1 is set for 9
Bit at index 4 and 1 is set for 11
Bit at index 13 and 5 is set for 13
Bit at index 5 and 6 is set for 15
Bit at index 15 and 0 is set for 17
Bit at index 1 and 0 is set for 19
2 is definitely not in the set.
4 is definitely not in the set.
6 may be in the set.
8 is definitely not in the set.
10 may be in the set.
12 is definitely not in the set.
14 is definitely not in the set.
16 is definitely not in the set.
18 is definitely not in the set.
```

## QUESTION 2:

Link Matrix:

```
[[0.  1.  0.  1. ]
 [0.  0.  0.5 0. ]
 [0.5 0.  0.  0. ]
 [0.5 0.  0.5 0. ]]
```

Pages: ['A', 'B', 'X', 'Y']

```
In [15]: import numpy as np
import networkx as nx
import matplotlib.pyplot as plt

def calculate_pagerank(connectivity_matrix, damping_factor=0.85, tolerance=1e-6):
    num_pages = connectivity_matrix.shape[1]
    pagerank_vector = np.ones((num_pages, 1)) / num_pages
    previous_pagerank_vector = np.ones((num_pages, 1)) * 100 # Initialize with a large difference
    pagerank_iterations = [] # Store PageRank vectors at each iteration
    while np.linalg.norm(pagerank_vector - previous_pagerank_vector, 2) > tolerance:
        previous_pagerank_vector = pagerank_vector.copy()
        pagerank_vector = damping_factor * np.dot(connectivity_matrix, pagerank_vector) + (1 - damping_factor) / num_pages
        pagerank_vector += ((1 - np.sum(pagerank_vector)) / num_pages) # Adding missing probability mass to handle dead-end nodes
        pagerank_vector = pagerank_vector / np.sum(pagerank_vector) # Normalize the PageRank vector
        pagerank_iterations.append(pagerank_vector.flatten())
    return pagerank_iterations

# input matrix representing the connectivity of web pages
webpage_connectivity = np.array([[0, 1, 0, 1],
                                  [0, 0, 0.5, 0],
                                  [0.5, 0, 0, 0],
                                  [0.5, 0, 0.5, 0]])

# Running the PageRank algorithm until convergence on the input matrix
page_rank_iterations = calculate_pagerank(webpage_connectivity)

# Create a directed graph using NetworkX
graph = nx.DiGraph()

# Add edges from the matrix
edges = [("A", "Y"), ("A", "X"), ("B", "A"), ("X", "B"), ("X", "Y"), ("Y", "A")]
graph.add_edges_from(edges)


# Draw the graph
positions = nx.spring_layout(graph)
nx.draw(graph, positions, with_labels=True, node_size=1500, node_color="skyblue", font_size=15, font_weight="bold", arrows=True)
plt.title("Graph Visualization")
plt.show()

# Print all iterations until convergence
for i, page_rank in enumerate(page_rank_iterations):
    print("Iteration", i+1, "Page ranks:", page_rank)

# Print the final iteration in the respective page rank order
final_page_ranks = page_rank_iterations[-1]
page_ranks_with_labels = [(page, rank) for page, rank in zip(['A', 'B', 'X', 'Y'], final_page_ranks)]
page_ranks_with_labels.sort(key=lambda x: x[1], reverse=True)

print("Final Page Ranks:")
for page, rank in page_ranks_with_labels:
    print(f"{page}: {rank}")
```

### Graph Visualization



```
Iteration 1 Page ranks: [0.4625 0.14375 0.14375 0.25 ]
Iteration 2 Page ranks: [0.3721875 0.09859375 0.2340625 0.29515625]
Iteration 3 Page ranks: [0.3721875 0.13697656 0.19567969 0.29515625]
Iteration 4 Page ranks: [0.40481289 0.12066387 0.19567969 0.27884355]
Iteration 5 Page ranks: [0.37708131 0.12066387 0.20954548 0.29270935]
Iteration 6 Page ranks: [0.38886723 0.12655683 0.19775956 0.28681638]
Iteration 7 Page ranks: [0.38886723 0.12154781 0.20276857 0.28681638]
Iteration 8 Page ranks: [0.38460957 0.12367664 0.20276857 0.28894522]
Iteration 9 Page ranks: [0.38822858 0.12367664 0.20095907 0.28713571]
Iteration 10 Page ranks: [0.3866905 0.1229076 0.20249715 0.28790475]
Iteration 11 Page ranks: [0.3866905 0.12356129 0.20184346 0.28790475]
Iteration 12 Page ranks: [0.38724613 0.12328347 0.20184346 0.28762693]
Iteration 13 Page ranks: [0.38677384 0.12328347 0.20207961 0.28786308]
Iteration 14 Page ranks: [0.38697457 0.12338383 0.20187888 0.28776272]
Iteration 15 Page ranks: [0.38697457 0.12329853 0.20196419 0.28776272]
Iteration 16 Page ranks: [0.38690206 0.12333478 0.20196419 0.28779897]
Iteration 17 Page ranks: [0.38696369 0.12333478 0.20193337 0.28776815]
Iteration 18 Page ranks: [0.3869375 0.12332168 0.20195957 0.28778125]
Iteration 19 Page ranks: [0.3869375 0.12333282 0.20194844 0.28778125]
Iteration 20 Page ranks: [0.38694696 0.12332809 0.20194844 0.28777652]
Iteration 21 Page ranks: [0.38693892 0.12332809 0.20195246 0.28778054]
Iteration 22 Page ranks: [0.38694233 0.12332979 0.20194904 0.28777883]
Iteration 23 Page ranks: [0.38694233 0.12332834 0.20195049 0.28777883]
Iteration 24 Page ranks: [0.3869411 0.12332896 0.20195049 0.28777945]
Iteration 25 Page ranks: [0.38694215 0.12332896 0.20194997 0.28777893]
Iteration 26 Page ranks: [0.3869417 0.12332874 0.20195041 0.28777915]
Final Page Ranks:
A: 0.3869417021334981
Y: 0.28777914893325096
X: 0.20195041300356312
B: 0.12332873592968788
```

### QUESTION 3

Link Matrix:

```
[ [0.      0.      0.5      0.33333333 0.      ]
  [0.      0.      0.5      0.33333333 0.      ]
  [0.      0.      0.      0.33333333 0.      ]
  [1.      0.      0.      0.      0.      ]
  [0.      1.      0.      0.      0.      ]]
```

Pages: ['A', 'Q', 'X', 'Y', 'Z']

```
In [16]: import numpy as np
import networkx as nx
import matplotlib.pyplot as plt

def calculate_pagerank(connectivity_matrix, damping_factor=0.85, tolerance=1e-6):
    num_pages = connectivity_matrix.shape[1]
    pagerank_vector = np.ones((num_pages, 1)) / num_pages
    previous_pagerank_vector = np.ones((num_pages, 1)) * 100 # Initialize with a large difference
    pagerank_iterations = [] # Store PageRank vectors at each iteration
    while np.linalg.norm(pagerank_vector - previous_pagerank_vector, 2) > tolerance:
        previous_pagerank_vector = pagerank_vector.copy()
        pagerank_vector = damping_factor * np.dot(connectivity_matrix, pagerank_vector) + (1 - damping_factor) / num_pages
        pagerank_vector += ((1 - np.sum(pagerank_vector)) / num_pages) # Adding missing probability mass to handle dead-end nodes
        pagerank_vector = pagerank_vector / np.sum(pagerank_vector) # Normalize the PageRank vector
        pagerank_iterations.append(pagerank_vector.flatten())
    return pagerank_iterations

# input matrix representing the connectivity of web pages
webpage_connectivity = np.array([[0.0,0.0,0.5,0.33333333,0.0],
 [0.0,0.0,0.5,0.33333333,0.0],
 [0.0,0.0,0.0,0.33333333,0.0],
 [1.0,0.0,0.0,0.0,0.0],
 [0.0,1.0,0.0,0.0,0.0]])

# Running the PageRank algorithm until convergence on the input matrix
page_rank_iterations = calculate_pagerank(webpage_connectivity)

# Create a directed graph using NetworkX
graph = nx.DiGraph()

# Add edges from the matrix
edges = [
    ('A', 'Y'),
    ('Y', 'A'),
    ('Y', 'X'),
    ('Y', 'Q'),
    ('X', 'A'),
    ('X', 'Q'),
    ('Q', 'Z')
]

graph.add_edges_from(edges)


# Draw the graph
positions = nx.spring_layout(graph)
nx.draw(graph, positions, with_labels=True, node_size=1500, node_color="skyblue", font_size=15, font_weight="bold", arrows=True)
plt.title("Graph Visualization")
plt.show()

# Print all iterations until convergence
for i, page_rank in enumerate(page_rank_iterations):
    print("Iteration", i+1, "Page ranks:", page_rank)

# Print the final iteration in the respective page rank order
final_page_ranks = page_rank_iterations[-1]
page_ranks_with_labels = [(page, rank) for page, rank in zip(['A','Q','X','Y','Z'], final_page_ranks)]
page_ranks_with_labels.sort(key=lambda x: x[1], reverse=True)

print("Final Page Ranks:")
for page, rank in page_ranks_with_labels:
    print(f"{page}: {rank}")
```

### Graph Visualization



```
Iteration 1 Page ranks: [0.20566667 0.20566667 0.12066667 0.234      0.234      ]
Iteration 2 Page ranks: [0.18736333 0.18736333 0.13608      0.24459667 0.24459667]
Iteration 3 Page ranks: [0.19871782 0.19871782 0.14088382 0.23084027 0.23084027]
Iteration 4 Page ranks: [0.19452321 0.19452321 0.13464759 0.23815299 0.23815299]
Iteration 5 Page ranks: [0.19518792 0.19518792 0.13796269 0.23583074 0.23583074]
Iteration 6 Page ranks: [0.19554408 0.19554408 0.13690993 0.23600095 0.23600095]
Iteration 7 Page ranks: [0.19517382 0.19517382 0.1369871  0.23633263 0.23633263]
Iteration 8 Page ranks: [0.19535698 0.19535698 0.13713746 0.2360743  0.2360743  ]
Iteration 9 Page ranks: [0.19530377 0.19530377 0.13702035 0.23618606 0.23618606]
Iteration 10 Page ranks: [0.19530466 0.19530466 0.13707101 0.23615983 0.23615983]
Iteration 11 Page ranks: [0.1953143  0.1953143  0.13705912 0.23615613 0.23615613]
Iteration 12 Page ranks: [0.19530757 0.19530757 0.13705745 0.2361637  0.2361637  ]
Iteration 13 Page ranks: [0.19531029 0.19531029 0.13706088 0.23615927 0.23615927]
Iteration 14 Page ranks: [0.19530974 0.19530974 0.13705887 0.23616082 0.23616082]
Iteration 15 Page ranks: [0.19530959 0.19530959 0.13705957 0.23616062 0.23616062]
Final Page Ranks:
Y: 0.23616062055337508
Z: 0.23616062055337508
A: 0.1953095925906417
Q: 0.1953095925906417
X: 0.13705957371196648
```

**Y=Z followed by A=Q and then X**

#### QUESTION 4

**Suppose we recursively eliminate dead ends from the graph, solve the remaining graph, and estimate the PageRank for the dead-end pages as described in Section 5.1.4. Suppose the graph is a chain of dead ends, headed by a node with a self-loop, as suggested in Fig. 5.9. What would be the Page Rank assigned to each of the nodes?**

The scenario outlined consists of a series of dead-end sites, where the initial node (Node 1) has a self-loop and each page links to the next. This is how each node's PageRank would be calculated:

**Node 1: Self-Looping Root Node** Node 1 keeps all of its significance and transfers its entire PageRank to itself because it has a self-loop. As a result, Node 1's PageRank stays at 1.

**Dead-End Nodes That Follow (Node 2 to Node n):** There is only one inbound link from each preceding node (Node 2 to Node n).

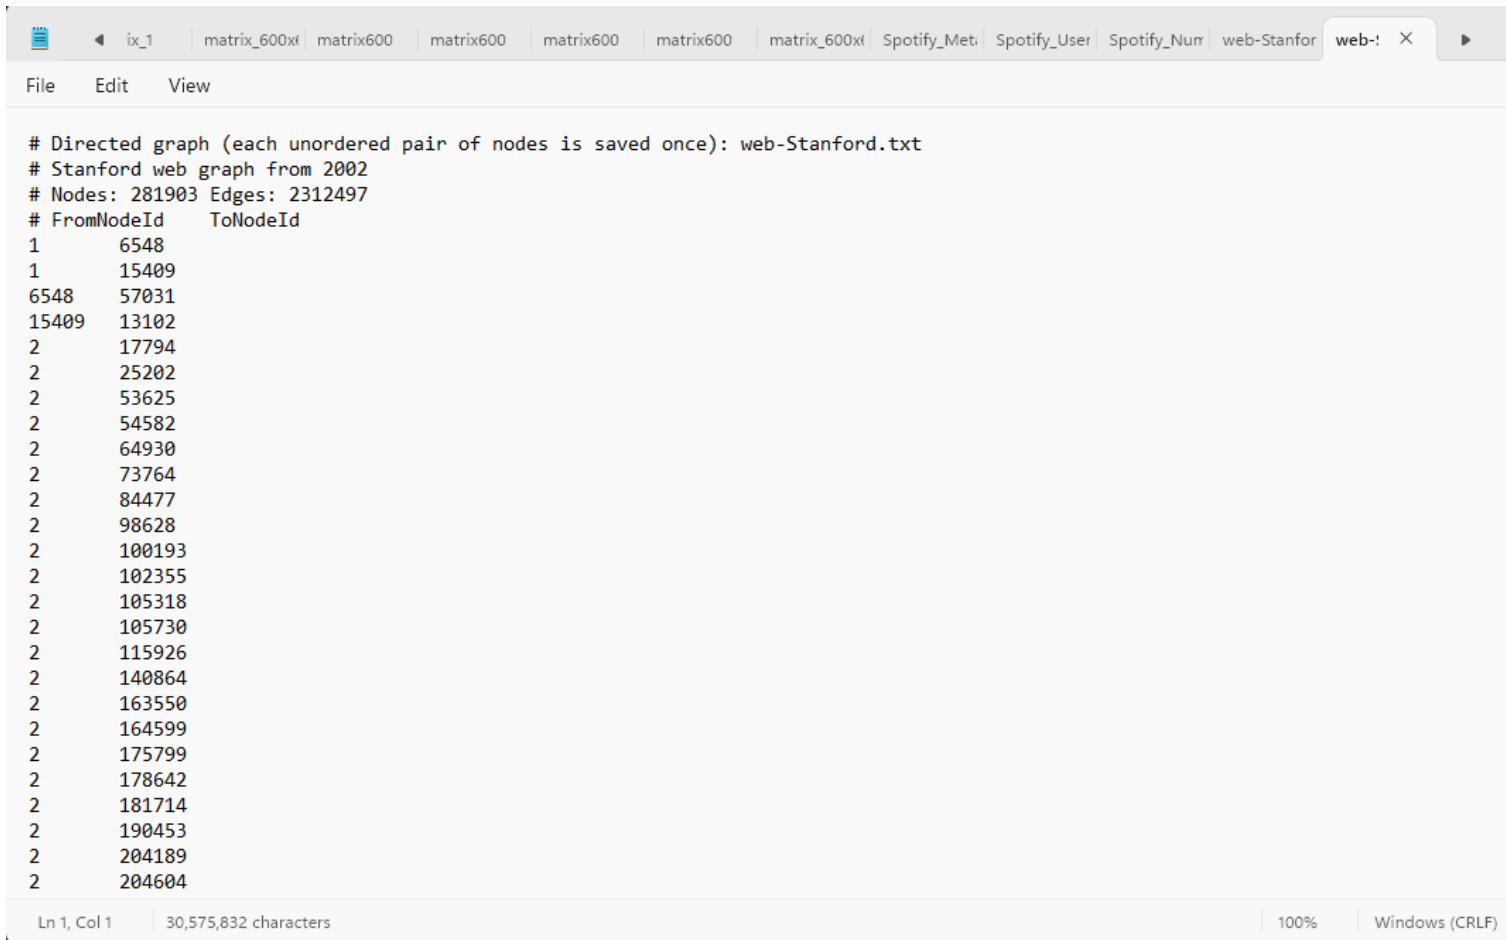
Hence, node 1, the root node with a self-loop, continues to have a PageRank of 1. Because each successive dead-end node (Node 2 to Node n) receives half of its PageRank from its predecessors, their PageRank is  $1/2$ . As a result, each node in the chain of dead ends, which is led by a node with a self-loop, would have the following PageRank assigned to it:

Node 1: 1

Node 2 to Node n:  $1/2$

## QUESTION 5:

- a. Take a look at the file and report how many nodes and edges the web-Stanford.txt contains.



```
# Directed graph (each unordered pair of nodes is saved once): web-Stanford.txt
# Stanford web graph from 2002
# Nodes: 281903 Edges: 2312497
# FromNodeId    ToNodeId
1        6548
1        15409
6548     57031
15409    13102
2        17794
2        25202
2        53625
2        54582
2        64930
2        73764
2        84477
2        98628
2        100193
2        102355
2        105318
2        105730
2        115926
2        140864
2        163550
2        164599
2        175799
2        178642
2        181714
2        190453
2        204189
2        204604
```

Ln 1, Col 1 | 30,575,832 characters | 100% | Windows (CRLF)

b. Report the runtime (took about 5 minutes to run when I tested it)

SINGLE NODE:

Real 1m47.097s

User 1m49.754s

Sys 0m3.330s

```
aws Services Search [Alt+S] N. Virginia voclabs/user3013569=paramah@depaul.edu @ 2854-6811-2301
24/03/10 00:44:44 INFO mapred.MapTask: Spilling map output
24/03/10 00:44:44 INFO mapred.MapTask: bufstart = 0; bufend = 4117449; bufvoid = 104857600
24/03/10 00:44:44 INFO mapred.MapTask: kstart = 26214396(104857684); kvend = 25086784(100347136); length = 1127613/6553600
24/03/10 00:44:44 INFO mapred.MapTask: Finished spill 0
24/03/10 00:44:44 INFO mapred.Task: Task:attempt_local1413088030_0010_m_000000_0 is done. And is in the process of committing
24/03/10 00:44:44 INFO mapred.LocalJobRunner: map
24/03/10 00:44:44 INFO mapred.Task: Task 'attempt_local1413088030_0010_m_000000_0' done.
24/03/10 00:44:44 INFO mapred.LocalJobRunner: Finishing task: attempt_local1413088030_0010_m_000000_0
24/03/10 00:44:44 INFO mapred.LocalJobRunner: map task executor complete.
24/03/10 00:44:44 INFO mapred.LocalJobRunner: Waiting for reduce tasks
24/03/10 00:44:44 INFO mapred.LocalJobRunner: Starting task: attempt_local1413088030_0010_r_000000_0
24/03/10 00:44:44 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
24/03/10 00:44:44 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@496033e8
24/03/10 00:44:44 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=334338464, maxSingleShuffleLimit=83584616, mergeThreshold=220663392, ioSortFactor=10, memToMemMergeOutputsThreshold=10
24/03/10 00:44:44 INFO reduce.EventFetcher: attempt_local1413088030_0010_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
24/03/10 00:44:44 INFO reduce.LocalFetcher: localFetcher#10 about to shuffle output of map attempt_local1413088030_0010_m_000000_0 decomp: 4681259 len: 4681263 to MEMORY
24/03/10 00:44:44 INFO reduce.inMemoryMapOutput: Read 4681259 bytes from map-output for attempt_local1413088030_0010_m_000000_0
24/03/10 00:44:44 INFO reduce.EventFetcher: closeInMemoryFile -> map-output of size: 4681259, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 4681259
24/03/10 00:44:44 INFO reduce.EventFetcher: EventFetcher is interrupted. Returning
24/03/10 00:44:44 INFO mapred.LocalJobRunner: 1 / 1 copied.
24/03/10 00:44:44 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
24/03/10 00:44:44 INFO mapred.Merger: Merging 1 sorted segments
24/03/10 00:44:44 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 4681249 bytes
24/03/10 00:44:44 INFO reduce.MergeManagerImpl: Merged 1 segments, 4681259 bytes to disk to satisfy reduce memory limit
24/03/10 00:44:44 INFO reduce.MergeManagerImpl: Merging 1 files, 4681263 bytes from disk
24/03/10 00:44:44 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
24/03/10 00:44:44 INFO mapred.Merger: Merging 1 sorted segments
24/03/10 00:44:44 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 4681249 bytes
24/03/10 00:44:44 INFO mapred.LocalJobRunner: 1 / 1 copied.
24/03/10 00:44:45 INFO mapreduce.Job: Job job_local1413088030_0010 running in uber mode : false
24/03/10 00:44:45 INFO mapreduce.Job: map 100% reduce 0%
24/03/10 00:44:45 INFO mapred.Task: Task:attempt_local1413088030_0010_r_000000_0 is done. And is in the process of committing
24/03/10 00:44:45 INFO mapred.LocalJobRunner: 1 / 1 copied.
24/03/10 00:44:45 INFO mapred.Task: Task attempt_local1413088030_0010_r_000000_0 is allowed to commit now
24/03/10 00:44:45 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1413088030_0010_r_000000_0' to hdfs://localhost/data/prOutput/result/_temporary/0/task_local1413088030_0010_r_000000
24/03/10 00:44:45 INFO mapred.LocalJobRunner: reduce > reduce
24/03/10 00:44:45 INFO mapred.Task: Task 'attempt_local1413088030_0010_r_000000_0' done.
24/03/10 00:44:45 INFO mapred.LocalJobRunner: Finishing task: attempt_local1413088030_0010_r_000000_0
24/03/10 00:44:45 INFO mapred.LocalJobRunner: reduce task executor complete.
24/03/10 00:44:46 INFO mapreduce.Job: map 100% reduce 100%
24/03/10 00:44:46 INFO mapreduce.Job: Job job_local1413088030_0010 completed successfully
24/03/10 00:44:46 INFO mapreduce.Job: Counters: 38
File System Counters
  FILE: Number of bytes read=4525537444
  FILE: Number of bytes written=4535392423
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=477365676
  HDFS: Number of bytes written=419097533
  HDFS: Number of read operations=269
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=112
Map-Reduce Framework
  Map input records=281904
  Map output records=281904
  Map output bytes=4117449
  Map output materialized bytes=4681263
  Input split bytes=115
  Combine input records=0
  Combine output records=0
  Reduce input groups=125886
  Reduce shuffle bytes=4681263
  Reduce input records=281904
  Reduce output records=281904
  Spilled Records=563808
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=8
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=908066816
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=22808184
File Output Format Counters
  Bytes Written=7508523
DONE!
real    1m47.097s
user    1m49.754s
sys     0m3.330s
[ec2-user@ip-172-31-62-135 src] 2024-03-10 00:44:46
$ []
i-062d65a3dd2e0d019 (nachiketh_server)
PublicIPs: 18.210.20.84 PrivateIPs: 172.31.62.135
```



aws

Services

Search

[Alt+S]

N. Virginia

vociabs/user3013569=nparamah@depauLedu @ 2854-6811-230

[ec2-user@ip-172-31-62-135 ~] 2024-03-10 00:55:31

\$ hadoop fs -cat /data/prOutput/result/part-r-00000 | more

0.0 154952

0.0 36711

0.0 154931

0.0 3671

0.0 36689

0.0 36702

0.0 36693

0.0 99979

0.0 1792

0.0 24032

0.0 179164

0.0 24033

0.0 262390

0.0 179163

0.0 100020

0.0 100022

0.0 9995

0.0 99949

0.0 17915

0.0 244555

0.0 262402

0.0 262403

0.0 99924

0.0 218029

0.0 262405

0.0 100048

0.0 194867

0.0 194854

0.0 218032

0.0 99904

0.0 100058

0.0 100059

0.0 100061

0.0 179206

0.0 244561

0.0 194850

0.0 100069

0.0 100074

0.0 218039

0.0 262422

0.0 198464

0.0 99869

0.0 100084

0.0 244567

0.0 100092

0.0 262338

0.0 262431

0.0 217985

0.0 100094

0.0 99841

0.0 10010

0.0 21798

0.0 262443

0.0 218050

0.0 100107

0.0 100109

0.0 244575

0.0 99823

0.0 100114

0.0 262449

0.0 217974

0.0 99814

0.0 262457

0.0 100124

0.0 217970

0.0 179104

0.0 262333

0.0 998

0.0 194826

0.0 99794

0.0 218063

0.0 218066

0.0 244580

0.0 99785

0.0 100137

0.0 198474

0.0 100148

0.0 99770

0.0 179093

0.0 262485

0.0 100158

0.0 99763

0.0 262329

0.0 100165

0.0 100168

0.0 100171

0.0 218073

0.0 9975

0.0 179085

0.0 217950

i-062d65a3dd2e0d019 (nachiketh\_server)

PublicIPs: 18.210.20.84 PrivateIPs: 172.31.62.135

## MULTI NODE:

Real 4m37.678S

User 0m6.563s

Sys 0m0.476s

```
aws Services Q Search [Alt+S] N. Virginia voclabs/user3013569=paramah@depaul.edu @ 2854-6811-230

GC time elapsed (ms)=226
CPU time spent (ms)=16110
Physical memory (bytes) snapshot=717959168
Virtual memory (bytes) snapshot=6141054976
Total committed heap usage (bytes)=603455488

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=22629528
File Output Format Counters
  Bytes Written=22629726

Running Job#3 (rank ordering) ...
24/03/10 01:11:58 INFO client.RMProxy: Connecting to ResourceManager at /172.31.57.221:8032
24/03/10 01:11:58 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
24/03/10 01:11:58 INFO input.FileInputFormat: Total input paths to process : 1
24/03/10 01:11:58 INFO mapreduce.JobSubmitter: number of splits:1
24/03/10 01:11:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1710032526505_0010
24/03/10 01:11:58 INFO impl.YarnClientImpl: Submitted application application_1710032526505_0010
24/03/10 01:11:58 INFO mapreduce.Job: The url to track the job: http://ip-172-31-57-221.ec2.internal:8088/proxy/application_1710032526505_0010/
24/03/10 01:11:58 INFO mapreduce.Job: Running job: job_1710032526505_0010
24/03/10 01:12:03 INFO mapreduce.Job: Job job_1710032526505_0010 running in uber mode : false
24/03/10 01:12:03 INFO mapreduce.Job:  map 0% reduce 0%
24/03/10 01:12:08 INFO mapreduce.Job:  map 100% reduce 0%
24/03/10 01:12:14 INFO mapreduce.Job:  map 100% reduce 100%
24/03/10 01:12:14 INFO mapreduce.Job: Job job_1710032526505_0010 completed successfully
24/03/10 01:12:14 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=4681263
  FILE: Number of bytes written=9577181
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=22629845
  HDFS: Number of bytes written=7333875
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=14448
  Total time spent by all reduces in occupied slots (ms)=12680
  Total time spent by all map tasks (ms)=3612
  Total time spent by all reduce tasks (ms)=3170
  Total vcore-milliseconds taken by all map tasks=3612
  Total vcore-milliseconds taken by all reduce tasks=3170
  Total megabyte-milliseconds taken by all map tasks=14448000
  Total megabyte-milliseconds taken by all reduce tasks=12680000

Map-Reduce Framework
  Map input records=281904
  Map output records=281904
  Map output bytes=4117449
  Map output materialized bytes=4681263
  Input split bytes=119
  Combine input records=0
  Combine output records=0
  Reduce input groups=123967
  Reduce shuffle bytes=4681263
  Reduce input records=281904
  Reduce output records=281904
  Spilled Records=563808
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=129
  CPU time spent (ms)=4090
  Physical memory (bytes) snapshot=508944394
  Virtual memory (bytes) snapshot=6136938496
  Total committed heap usage (bytes)=396361728

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=22629726
File Output Format Counters
  Bytes Written=7333875

DONE!

real    4m37.678s
user    0m6.563s
sys     0m0.476s
[ec2-user@ip-172-31-57-221 src]$ []

i-06fafa7fc816ec90a (Master_1)
PublicIPs: 54.90.40.172 PrivateIPs: 172.31.57.221
```

c. Submit a screenshot of the first page of nodes, e.g., by running:

```
aws Services Search [Alt+S] N. Virginia voclabs/user3013569=npamah@depaul.edu @ 2854-6811-230

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=22629726
File Output Format Counters
  Bytes Written=7333875
DONE!

real    4m37.678s
user    0m6.563s
sys      0m0.476s
[ec2-user@ip-172-31-57-221 src]$ hadoop fs -cat /data/prOutput/result/part-r-00000 | more
0.15000000596046448 75380
0.15000000596046448 155237
0.15000000596046448 75378
0.15000000596046448 155226
0.15000000596046448 155222
0.15000000596046448 155221
0.15000000596046448 75372
0.15000000596046448 125860
0.15000000596046448 47294
0.15000000596046448 155216
0.15000000596046448 155204
0.15000000596046448 155203
0.15000000596046448 47303
0.15000000596046448 155116
0.15000000596046448 125881
0.15000000596046448 47316
0.15000000596046448 75340
0.15000000596046448 47317
0.15000000596046448 125872
0.15000000596046448 75358
0.15000000596046448 75357
0.15000000596046448 47327
0.15000000596046448 47336
0.15000000596046448 155122
0.15000000596046448 47357
0.15000000596046448 155140
0.15000000596046448 47360
0.15000000596046448 155137
0.15000000596046448 75347
0.15000000596046448 155129
0.15000000596046448 68099
0.15000000596046448 229465
0.15000000596046448 68079
0.15000000596046448 100022
0.15000000596046448 22977
0.15000000596046448 229439
0.15000000596046448 68199
0.15000000596046448 229784
0.15000000596046448 9995
0.15000000596046448 99949
0.15000000596046448 229792
0.15000000596046448 68078
0.15000000596046448 229796
0.15000000596046448 22941
0.15000000596046448 1335
0.15000000596046448 229811
0.15000000596046448 229401
0.15000000596046448 229388
0.15000000596046448 229387
0.15000000596046448 229831
0.15000000596046448 68068
0.15000000596046448 229369
0.15000000596046448 229842
0.15000000596046448 229846
0.15000000596046448 22936
0.15000000596046448 229358
0.15000000596046448 133522
0.15000000596046448 229855
0.15000000596046448 229351
0.15000000596046448 133529
0.15000000596046448 229341
0.15000000596046448 99924
0.15000000596046448 229871
0.15000000596046448 100048
0.15000000596046448 229313
0.15000000596046448 229304
0.15000000596046448 229288
0.15000000596046448 229282
0.15000000596046448 229288
0.15000000596046448 229284
0.15000000596046448 229920
0.15000000596046448 229922
0.15000000596046448 100058
0.15000000596046448 100059
0.15000000596046448 229927

i-06fafa7fc816ec90a (Master_1)
PublicIPs: 54.90.40.172 PrivateIPs: 172.31.57.221
```